

**Sample Midterm**  
**ECE 454S 2014: Computer Systems Programming**  
**Date: Monday, October 27, 2014 7 p.m. - 9 p.m.**

Instructor: Cristiana Amza  
Department of Electrical and Computer Engineering  
University of Toronto

Problem number	Maximum Score	Your Score
1	7	
2	23	
3	20	
4	20	
5	27	
6	13	
total	100	

This exam is open textbook and open lecture notes. You have two hours to complete the exam. Use of computing and/or communicating devices is NOT permitted. You should not need any such devices. You can use a basic calculator if you feel it is absolutely necessary.

Do not remove any sheets from this test book. Answer all questions in the space provided. No additional sheets are permitted. Scratch space is available at the end of the exam. Work independently.

Write your name and student number in the space below. Do the same on the top of each sheet of this exam book.

Your Student Number:

Your First Name:

Your Last Name:

Student Number:

Name:

---

**Problem 1. Basic Code Optimization Facts. (7 Points)**

Name two optimization blockers and briefly explain what they are. (4 points)

A sorting algorithm takes 1 second to execute. What are the advantages and disadvantages of using gprof as a profiling tool in order to find out the bottleneck of this algorithm ? (3 points)

Student Number:

Name:

**Problem 2. Performance Optimization. (23 Points)**

The following problem concerns optimizing codes for maximum performance on an Intel Pentium III. Recall the following performance characteristics of the functional units for this machine:

Operation	Latency	Issue Time
Integer Add	1	1
Integer Multiply	4	1
Integer Divide	36	36
Floating Point Add	3	1
Floating Point Multiply	5	2
Floating Point Divide	38	38
Load or Store (Cache Hit)	1	1

**Part A**

Consider the following code segments:

Loop 1	Loop 2
for (i = 0; i < n; i++) x = y * a[i];	for (i = 0; i < n; i++) x = x * a[i];

When compiled with GCC, we obtain the following assembly code for the loop:

Loop 1	Loop 2
.L21: movl %ecx,%eax imull (%esi,%edx,4),%eax incl %edx cmpl %ebx,%edx jl .L21	.L27: imull (%esi,%edx,4),%eax incl %edx cmpl %ebx,%edx jl .L27

Running on one of the cluster machines, we find that Loop 1 requires 3.0 clock cycles per iteration, while Loop 2 requires 4.0.

- (a) (6 points) Explain how it is that Loop 1 is faster than Loop 2, even though it has one more instruction.

Student Number:

Name:

---

(b) (6 points) We perform 4-way loop unrolling for the two loops. This speeds up Loop 1. Briefly explain why.

(c) (6 points) Even with loop unrolling, we find that the performance of Loop 2 remains the same. Briefly explain why.

## Part B

You've just joined a programming team that is trying to develop the world's fastest factorial routine. Starting with recursive factorial, they've converted the code to use iterations as follows:

```
int fact(int n) {
    int i;
    int result = 1;

    for (i = n; i > 0; i--)
        result = result * i;

    return result;
}
```

By doing so, they have reduced the number of cycles per element (CPE) for the function from around 63 to around 4 (really!). Still, they would like to do better.

Student Number:

Name:

---

One of the programmers heard about loop unrolling. He generated the following code:

```
int fact_u2(int n) {
    int i;
    int result = 1;

    for (i = n; i > 1; i-=2) {
        result = (result * i) * (i-1);
    }

    return result;
}
```

(a) (4 points) However, benchmarking `fact_u2` shows no improvement in performance. How would you explain that?

(b) (5 points) You modify the line inside the loop to read:

```
result = result * (i * (i-1));
```

To everyone's astonishment, the measured performance now has a CPE of 2.5. How do you explain this performance improvement?

Student Number:

Name:

---

### Problem 3. Cache Miss Rate. (20 Points)

You are writing a new 3D game that you hope will earn you fame and fortune. You are currently working on a function to blank the screen buffer before drawing the next frame. The screen you are working with is a 640x480 array of pixels. The machine you are working on has a 64 KB direct mapped cache with 4 byte lines. The C structures you are using are:

```
struct pixel {
    char r;
    char g;
    char b;
    char a;
};

struct pixel buffer[480][640];
register int i, j;
register char *cptr;
register int *iptr;
```

Assume:

- `sizeof(char) = 1`
- `sizeof(int) = 4`
- `buffer` begins at memory address 0
- The cache is initially empty.
- The only memory accesses are to the entries of the array `buffer`. Variables `i`, `j`, `cptr`, and `iptr` are stored in registers.

Student Number:

Name:

---

A. (5 points) What percentage of the writes in the following code will miss in the cache?

```
for (j=0; j < 640; j++) {  
    for (i=0; i < 480; i++){  
        buffer[i][j].r = 0;  
        buffer[i][j].g = 0;  
        buffer[i][j].b = 0;  
        buffer[i][j].a = 0;  
    }  
}
```

Miss rate for writes to buffer: \_\_\_\_\_ %

B. (5 points) What percentage of the writes in the following code will miss in the cache?

```
char *cptr;  
cptr = (char *) buffer;  
for (; cptr < (((char *) buffer) + 640 * 480 * 4); cptr++)  
    *cptr = 0;
```

Miss rate for writes to buffer: \_\_\_\_\_ %

C. (5 points) What percentage of the writes in the following code will miss in the cache?

```
int *iptr;  
iptr = (int *) buffer;  
for (; iptr < (buffer + 640 * 480); iptr++)  
    *iptr = 0;
```

Miss rate for writes to buffer: \_\_\_\_\_ %

D. (5 points) Which code (A, B, or C) should be the fastest? \_\_\_\_\_

Student Number:

Name:

#### Problem 4 . Cache Conflict Misses. (20 Points)

Consider the following matrix transpose routine

```
typedef int array[2][2];

void transpose(array dst, array src) {
    int i, j;

    for (i = 0; i < 2; i++) {
        for (j = 0; j < 2; j++) {
            dst[j][i] = src[i][j];
        }
    }
}
```

running on a hypothetical machine with the following properties:

- `sizeof(int) == 4`.
- The `src` array starts at address 0 and the `dst` array starts at address 16 (decimal).
- There is a single L1 cache that is direct mapped with a cache line size of 8 bytes. The address mapping to the cache line is as follows:  $\text{cache line number} = (\text{address} \% \text{cache size}) / \text{block size}$ .
- Accesses to the `src` and `dst` arrays are the only sources of read and write misses, respectively.

#### Part A (6 points)

Suppose the cache has a total size of 16 data bytes (i.e., the block size times the number of sets is 16 bytes) and that the cache is initially empty. Then for each `row` and `col`, indicate whether each access to `src[row][col]` and `dst[row][col]` is a hit (h) or a miss (m). For example, reading `src[0][0]` is a miss and writing `dst[0][0]` is also a miss.

dst array	
dst[0][0]	m
dst[1][0]	
dst[0][1]	
dst[1][1]	

src array	
src[0][0]	m
src[0][1]	
src[1][0]	
src[1][1]	



Student Number:

Name:

---

**Part B (6 points)**

Repeat part A for a cache with a total size of 32 data bytes.

dst array	
dst[0][0]	m
dst[1][0]	
dst[0][1]	
dst[1][1]	

src array	
src[0][0]	m
src[0][1]	
src[1][0]	
src[1][1]	


Student Number:

Name:

### Problem 5. Dynamic Memory Allocation. (27 Points)

#### Part A. (17 points)

Consider a memory allocator that uses an implicit free list. The layout of each allocated and free memory block is as follows:



Each memory block, either allocated or free, has a size that is a multiple of eight bytes. Thus, only the 29 higher order bits in the header and footer are needed to record block size, which includes the header and footer. The usage of the remaining 3 lower order bits is as follows:

- bit 0 indicates the use of the current block: 1 for allocated, 0 for free.
- bit 1 indicates the use of the previous adjacent block: 1 for allocated, 0 for free.
- bit 2 is unused and is always set to be 0.

Note: The header and footer will always be present regardless of whether the block is allocated or not.

Student Number:

Name:

Given the contents of the heap shown on the left, show the new contents of the heap (in the right table) after a call to `free(0x400b010)` is executed. Your answers should be given as hex values. Note that the address grows from bottom up. Assume that the allocator uses immediate coalescing, that is, adjacent free blocks are merged immediately each time a block is freed. Please also briefly explain your answer.

Address

0x400b028	0x00000012
0x400b024	0x400b611c
0x400b020	0x400b512c
0x400b01c	0x00000012
0x400b018	0x00000013
0x400b014	0x400b511c
0x400b010	0x400b601c
0x400b00c	0x00000013
0x400b008	0x00000013
0x400b004	0x400b601c
0x400b000	0x400b511c
0x400affc	0x00000013

Address

0x400b028	
0x400b024	0x400b611c
0x400b020	0x400b512c
0x400b01c	
0x400b018	
0x400b014	0x400b511c
0x400b010	0x400b601c
0x400b00c	
0x400b008	
0x400b004	0x400b601c
0x400b000	0x400b511c
0x400affc	

**Explanation:** (You can use the space in the next page):

Student Number:

Name:

---

Student Number:

Name:

---

### Part B. 10 Points

Assume that you want to extend the previous implicit allocator to improve its performance. You would like to reduce allocation time by maintaining an explicit doubly-linked free list. In addition, you would like to improve its memory utilization by using the footer only when a block is free. You decide that a first-fit search algorithm is sufficient. You may assume that: `sizeof(void *) == 4`.

(a) Given that the block size must be a multiple of 8 bytes, what is the minimum block size allowable under this scheme?

(b) Determine the amount of memory (in bytes), wasted due to internal fragmentation, after the following four allocation requests. Do not include the cost of memory used for block headers in your count.

`malloc(1)`

`malloc(5)`

`malloc(12)`

`malloc(13)`

In order to further improve the performance of your allocator, you decide to try to implement an explicit binary tree data structure to enable a fast best-fit search through the free blocks. Each free block within the tree must now maintain a pointer to each of its children, and to its parent.

(c) Assuming that the block size must still be a multiple of 8 bytes, what is the minimum block size allowable under this new scheme?

(d) Comment on the effect that this allocator has on memory utilization when compared to the previous explicit linked list first-fit allocator. You should discuss any opposing tensions that might exist.

Student Number:

Name:

---

## Problem 6. Profiling and Speedup. (13 Points)

### Part A (5 points)

Explain the following profile data obtained using gprof on the n-body simulation application. Note: You have to explain the meaning of the numbers and the organization of the output. What conclusions can you draw from this profile data? Think of why we used gprof in the first place when drawing conclusions.

Flat profile sample: .....

% time	cumulative seconds	self seconds	calls	self s/call	total s/call	name
96.33	1076.18	1076.18	278456	0.00	0.00	World::update()
0.96	1086.85	10.67	212849617	0.00	0.00	ThreadManagerSerial::doSerialWork()
0.63	1093.86	7.01	1	7.01	39.30	OpenGLWindow::runWindowLoop (char const*, int, int, bool (*)(()), void (*)(SDL_KeyboardEvent*))

.....

Call graph sample:

index	% time	self	children	called	name
.....		1076.18	0.09	278456/278456	worldUpdate() [5]
[6]	96.3	1076.18	0.09	278456	World::update() [6]
		0.07	0.01	278240/278241	World::calculateTimeDeltaMilliseconds()[2]
		0.01	0.00	5338/5538	
					World::resetForegroundElement(World::ForegroundElement*) [38]
		0.00	0.00	149/149	World::updateCallsPerSecond(int) [47]
		0.00	0.00	4052/4287	World::radiusForMass(float) [139]
		0.00	0.00	642/642	
					World::swapForegroundElement(World::ForegroundElement*,World::ForegroundElement*)[140]

.....

Please also use the space on the next page

Student Number:

Name:

---

**Part B (8 Points)**

Assume that by profiling the code of our game application with gprof you get that 80% of the time is spent in a loop which is updating the physics of the game world, while 20% of the time is spent in frame rendering. Therefore you focus all your efforts only on improving the update world function through code optimizations such as loop unrolling, use of blocking for the update loop for cache hit improvements and loop parallelization. What is the best speedup you can get for the game application as a whole ?

Student Number:

Name:

---